

13 August 2012

3D-ICE 2.2.2

This document provides a brief summary of the usage of 3D-ICE 2.2.2. This includes illustrative examples of generating the input stack and floorplan files, and the various functions used for printing the results.

User Guide

Table of Contents

1. License and Copyright	3
2. What is new in 2.x?	4
3. Who needs 3D-ICE?.....	5
4. Before you begin.....	6
A. Compile SuperLU.....	6
B. Compile 3D-ICE	6
5. Overview of 3D-ICE	8
A. Principle of thermal simulation.....	8
i. Microchannel 4-resistor model.....	8
ii. Microchannel 2-resistor model.....	9
iii. Pin fins in-line.....	9
iv. Pinfins staggered.....	10
B. Inputs to 3D-ICE	11
C. Convention and Terminology.....	11
6. Creating a 3D-ICE project	13
A. Stack Description File.....	13
i. Materials	13
ii. Layers	14
iii. Dies.....	14
iv. Heat Sink	15
v. Liquid-cooled cavity	16
vi. Stack.....	20
vii. Dimensions.....	22
viii. Analysis options	24
ix. Output Instructions.....	25
Examples	28
B. Floorplan File.....	30
Time Slots.....	32
7. Network interface for remote simulations	34
8. Running 3D-ICE	35

9.	Usage of the 3D-ICE as Software Thermal Library	36
A.	StackDescription_t, Analysis_t and Output_t.....	37
B.	ThermalData_t	37
C.	Emulation and thermal output	38
D.	Socket	38
E.	NetworkMessage	39
F.	Debugging of Thermal Simulation.....	41
10.	References	42

1. License and Copyright

This file is part of 3D-ICE 2.2.2.

3D-ICE is free software: you can redistribute it and/or modify it under the terms of the GNU General Public License as published by the Free Software Foundation, either version 3 of the License, or any later version. 3D-ICE is distributed in the hope that it will be useful, but WITHOUT ANY WARRANTY; without even the implied warranty of MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU General Public License for more details. You should have received a copy of the GNU General Public License along with 3D-ICE. If not, see <http://www.gnu.org/licenses/>.

Any usage of 3D-ICE for research, commercial or other purposes must be properly acknowledged in the resulting products or publications. Specifically, [1] and [2] must be cited in these cases.



Copyright©2010,

Embedded Systems Laboratory -École Polytechnique Fédérale de Lausanne,

All Rights Reserved.

Authors:

*Arvind Sridhar**
*Alessandro Vincenzi**
*Giseong Bak**
*Martino Ruggiero**
Thomas Brunschwiler†
*David Atienza**



Contact Information:

EPFL-STI-IEL-ESL
Bâtiment ELG, ELG 130
Station 11
1015 Lausanne, Switzerland

Email: 3d-ice@listes.epfl.ch (SUBSCRIPTION NECESSARY!)
URL: <http://esl.epfl.ch/3d-ice.html>

This research has been partially funded by the Nano-Tera RTD project CMOSAIC (ref.123618) - which is financed by the Swiss Confederation and scientifically evaluated by SNSF, and the PRO3D project- financed by the European Community 7th Framework Programme (ref.FP7-ICT-248776).

* Embedded Systems Laboratory, Department of Electrical Engineering, EPFL, Lausanne, Switzerland.

† Advanced Thermal Packaging Group, IBM Research Laboratory, Zurich, Switzerland.

2. What is new in 2.x?

3D-ICE 2.x features several new additions and improvements over the first release of the simulator in September 2010. You would find the following main changes in this package (changes specific to version 2.2 are highlighted in **bold**):


- A new porous medium model has been implemented for the simulation of convective heat transport in interlayer microchannel cavities.
- Various new and enhanced heat transfer geometries have been introduced for the interlayer cavities.
- The user interface has been improved considerably. All the input-output information to the simulator has been transferred to the stack descriptor file, requiring little or no programming of the main C file by the user. This feature brings it closer to SPICE-like circuit simulators that use netlists for circuit descriptions. Hence, 3D-ICE can now function as both a stand-alone thermal simulator as well as a software thermal library. This functionality allows generating thermal **and power** maps of the IC for each die of the 3D stack.
- The user now has the option of doing steady state thermal analysis in addition to transient analysis. This is useful for obtaining steady state temperatures for some average/corner case heat dissipation scenarios, without having to wait for the transients to settle; as well as for generating initial temperature states for certain types of transient simulations.
- A new network interface has been created for online thermal simulations of IC architectures, in tandem with functional emulation on other devices. Under this system, 3D-ICE installed on a computer acts as the server, which communicates with a device emulating the desired architecture via an Ethernet cable. This feature is particularly useful for those who wish to build and test online-thermal-performance-management schemes with Hardware-Software co-simulation of ICs with liquid cooling without actually building the device [6].
- Extensive new documentation has been provided with this release to help simplify the usage of the simulator— this includes this extended user guide, research publications and a new *Doxygen* source code documentation. Doxygen has been used to create an organized html-based visualization of the 3D-ICE software library files and can be accessed by running in the `make doc` command inside the 3D-ICE folder, and then opening `3d-ice/doc/html/index.html` on your web browser.
- Floorplan elements can now be of irregular shapes, as long as they can be represented as a combination of rectangles (defined as an IC element). This feature has been introduced to support the more complex floorplanning that is seen in realistic architectures. **A new technique to align the IC elements on the underlying thermal cells has been implemented. It allows declaring IC elements smaller than thermal cells as well as posing more than one IC element on the same thermal cell (see section 6.B).**
- **The heat sink model has been extended to support two different techniques to model the dissipation of the heat through the top surface of the IC (see section 6.A.IV).**

To discover the **changes in the library that are related to the implementation and the interface of the software library** please see the CHANGELOG file in the package.

3. Who needs 3D-ICE?

3D-ICE, or “**3D Interlayer Cooling Emulator**”, is a Linux based *generic simulation platform* written in C to simulate the transient thermal behavior of 3D IC structures with inter-tier microchannel heat sinks. It is intended for various purposes including, but not limited to:

- Performing thermal analysis of 2D or 3D ICs during early stages of VLSI circuit/architecture design by electronic engineers.
- Simulating run-time and design-time thermal management strategies such as DVFS, dynamic work allocation, variable coolant flow rate, architectural floorplanning etc.
- Testing of microchannel heat-sink performances by microfabrication engineers and heat-sink designers.
- Evaluating accuracies of new and existing heat transfer correlations by experimental heat transfer engineers.



3D-ICE is based on the conventional compact modeling of heat transfer by conduction in solids, and advances a novel compact modeling methodology, called the Compact Transient Thermal Modeling (CTTM), for heat transfer by convection in microchannels. The user is free to use microchannel heat sinks of any dimension with the corresponding heat transfer performance data depending upon the accuracy/speed needs of the user. This simulator is ideal for situations where a quick estimate of chip temperatures is required, when the electronic designer is still iterating between various floorplanning and operating strategies in order to optimize for electronic performance and thermal safety/reliability of the final system.

In addition, the format of inputs, outputs and the problem construction/solving in 3D-ICE have been modeled on the popular compact modeling simulator **HotSpot**, making it easier for users who are familiar with this tool or with conventional circuit simulators. With numerous functions to access a variety of thermal data during the simulation, the user can reach deep into the heart of the thermal simulation, use the data to interface with other (popular or custom) tools and automate any kind of design/run-time optimization algorithms using 3D-ICE. More functionality will be added in the future versions of 3D-ICE to make this interfacing even more automatic and easy.

For more details on the theory and discussions about the accuracy/speed of the modeling technique used in 3D-ICE, please refer to the publications[1] and [2] available with this library.

4. Before you begin

The 3D-ICE library has been written and developed using:

- bison 2.4.1
- flex 2.5.35
- gcc 4.1.2

Make sure that these tools are installed on your system before compiling and that the corresponding variables in `makefile.def` point to the respective binary file. To use 3D-ICE, you must also download the SuperLU 4.3 library [4], available at <http://crd.lbl.gov/~xiaoye/SuperLU/>.

A. Compile SuperLU

Before compiling 3D-ICE, you must compile the SuperLU library by executing the following commands:

```
$ wget http://crd.lbl.gov/~xiaoye/SuperLU/superlu_4.3.tar.gz
$ tar xvfz superlu_4.3.tar.gz
$ cd SuperLU_4.3/
$ cp MAKE_INC/make.linux make.inc
```

Next, check and edit the `SuperLUroot` variable in `./make.inc` and select the `blas` library before compiling. You can either use a `blas` library installed on your system or the `blas` library supplied by the authors of SuperLU (see the README file). If you decide to use the former then the variable `BLASDEF` must be set to `-DUSE_VENDOR_BLAS` and `BLASLIB` must point to your `blas` library. Then compile SuperLU with

```
$ make superlulib
```

For the latter, `BLASDEF` must be unset and `BLASLIB` must point to the `libblas.a` archive. To compile the `blas` library included in the sources of SuperLU the Fortran compiler `g77` must be available on the system. Then compile SuperLU with

```
$ make blaslib
$ make superlulib
```

These are the operations that can be done when compiling SuperLU on a generic Linux platform. In case of a different architecture, please reference to the README file.

B. Compile 3D-ICE

Check and edit the `SLU_VERSION` and `SLU_MAIN` variables in `./makefile.def` to make it point to the main folder of SuperLU. Next, select the value of `SLU_LIBS` according to the choice done above when compiling SuperLU. You can then compile 3D-ICE with:

```
$ tar xvfz 3d-ice-2.2.2.tar.gz
$ cd 3d-ice
```

```
$ make
```

Now you can find the executable `3D-ICE-Emulator` in the `3d-ice/bin/` folder: this binary file will serve as the thermal simulator application for all your 3D-ICE projects. The `bin` folder will also contain two other binaries `3D-ICE-Server` and `3D-ICE-Client`: these two executables will serve as examples of an implementation of the network interface to execute a remote client/server thermal simulation.

5. Overview of 3D-ICE

This chapter provides a general overview of 3D-ICE: its operating principle, the input files required and the conventions and terminology used in writing these input files. In the next chapter, these input files are discussed in detail.

A. Principle of thermal simulation

3D-ICE is based on the compact modeling of heat flow in solids and liquids applied to a 3D-IC structure with microchannel cooling. As quick recap, the structure is divided into cuboidal *thermal cells* based on the discretization parameters you provide. Next, thermal conductance for heat flow through each face of the cuboid is calculated and connected to the neighboring cells at these faces. Also, a capacitance representing the heat capacity of the cell is calculated. Hence, an equivalent electrical RC circuit is created where the temperatures are represented by voltages and the heat flow is represented by the currents in the circuit. A typical thermal cell representing transient conduction in solids is shown below:




Figure 1: A typical solid thermal cell

In the case of thermal cells representing microchannels or fluidic cavities, while the heat transferred from the cavity walls into the fluid is represented using conventional thermal resistances (or conductances), the convective heat transport in the downstream direction due to mass flow is represented using voltage controlled current sources. There are four different interlayer cooling cavity models implemented in 3D-ICE.

i. Microchannel 4-resistor model

This is the model, which was available in 3D-ICE 1.0, is based on the 4RM-CTTM model described in [1]. Here, the thermal cells are constructed such that in the cavity layer, 2 faces of the thermal cells corresponding to the microchannel completely cover the entire cross-section of the microchannel as shown in Fig. 2. Hence, the discretization along the direction perpendicular to the channels is fixed by the channel geometry. There are four resistances one for each wall of the microchannel representing convective cooling of the wall surfaces. They can be coupled with solid thermal cells on all four directions as shown in Fig. 2. The convective resistance in the top, bottom and the sides can be provided independently into the simulator.




Figure 2: A mc4rm fluid thermal cell adjacent to two solid thermal cells

ii. Microchannel 2-resistor model

This new model for microchannel cavities is based on the 2RM-CTTM model presented in [2]. Here, the dependence of discretization on the channel geometry is completely removed, by homogenizing the entire cavity layer into a single “porous” material. Hence, each cell consists of circuit parameters corresponding to both the flowing coolant as well as the solid walls as shown in Fig. 3. The thermal cell here is two-dimensional because of the elimination of convective resistances on the sides. This model lets the user completely control the granularity of the model by making the thermal cell dimensions (in the x-direction) completely independent of the channel dimensions. Also, when using this model, the top and bottom heat transfer coefficients to be provided in the input files of 3D-ICE must be a projected average of the heat transfer coefficients between the side walls, and the top and bottom walls respectively. For more information about constructing this model, refer to [2].




Figure 3: A mc2rm thermal cell

iii. Pinfins in-line

Enhanced heat transfer geometries, such as pin fins can be used in interlayer cooling to improve the heat transfer properties of the liquid-cooled heat sink, when compared to conventional microchannels. Pin fins are also attractive because they can serve as ideal paths for the fabrication of TSVs in 3D ICs. These advantages come at the cost of increased pressure drops. Inline pin fins are one of the two standard types of pin fin HTGs. The top view of such an HTG is shown in Fig. 4.




Figure 4: Pin fins inline heat transfer geometry

iv. Pinfins staggered

Staggered pin fins (Fig. 5) is another standard type of pin fin heat transfer geometry. For both these pin fin geometries the 2RM-CTTM modeling method is utilized. The corresponding “porous” thermal cell is illustrated in Fig. 6.




Figure 5: Pin fins staggered heat transfer geometry




Figure 6: A pin fin thermal cell

In 3D-ICE, the heat transfer coefficients calculations for two specific test cases have been implemented that uses the darcy velocity of the coolant in the cavity. These computations for the convective resistance are based on empirical studies on these geometries performed in [3] and the resulting formulae (where v_{darcy} is the darcy velocity in m/s):

$$\begin{aligned}
\text{HTC}_{\text{inline}} &= \frac{2.526 \times 10^{-5}}{\left(\frac{v_{\text{darcy}}}{1 \text{ m/s}} + 1.35 \right)^{0.64}} + 1.533 \times 10^{-6} \\
\text{HTC}_{\text{staggered}} &= \frac{2.526 \times 10^{-5}}{\left(\frac{v_{\text{darcy}}}{1 \text{ m/s}} + 1.35 \right)^{1.52}} + 1.533 \times 10^{-6}
\end{aligned}$$

B. Inputs to 3D-ICE

3D-ICE accesses all the information needed to emulate a 3D IC from two different types of input files and these constitute a 3D-ICE project:

- **Stack Description File:** The stack description file or the “stack descriptor” is the project file created by the user to describe the 3D IC thermal problem that will be solved. It contains information about the structure, material properties of the 3D Stack, the description of the various heat sinks in the system, the discretization parameters, analysis parameters, and finally, commands to 3D-ICE for printing out the desired outputs from the simulation. Please refer to Section 6.A when writing this file.
- **Floorplan File:** This file describes the architecture of a 3D IC design as is seen by the thermal model. It contains information about the location and the size of each major logic block in the IC, and the corresponding heat dissipation traces, as a function of time. Each die included in the stack descriptor must have a corresponding floorplan file. Please refer to Section 6.B for instructions on how to write this file.

C. Convention and Terminology

In 3D-ICE, both Cartesian coordinates and cardinal directions are used to describe the location of cells/nodes and direction of heat flow in the structure. The indices of cells/nodes along the x direction (WEST-EAST) are sometimes referred to as *columns*, the indices along the y direction (SOUTH-NORTH) are sometimes referred to as *rows*, and the indices along the z direction (BOTTOM-TOP) sometimes referred to as *layers*. Also, the word *length* is used primarily to refer to dimensions in the x direction, the term *width* is used for the y direction and the term *height* is used for the z direction unless otherwise specified. A typical solid thermal cell along with the coordinate systems used in the library is shown in Fig. 1.

The corresponding model for a liquid cell is shown in Fig. 2. Note that the current sources shown here correspond to the fluid flow in the microchannels. In this library, only flow direction NORTH or +y is supported, and hence, when microchannels used in a 3D-IC structure, they are always laid out facing SOUTH-NORTH (with the inlet being at the southern end and the outlet at the northern end of the channels). Hence, the northern edge of the IC is expected to be the hottest and the southern end of the IC is expected to be the coldest in any analysis. It is up to you to decide your floorplanning of the ICs accordingly.

The syntaxes used in this document for describing how these input files must be written are based on the following convention:

<code>[: ... :]</code>	POSIX characters class
<code>... ...</code>	OR- either of the two elements must be used
<code>[...]?</code>	an optional element
<code>[...]+</code>	one or more of this element must be used
<code>[...]*</code>	zero or more of this element must be used

Within the files, **keywords must be written in low case** and all the white spaces belonging to

`[:space:]`

will be skipped during the parsing. Identifiers (referred to as ID) must match the following expression:

`[:alpha:] [_ | [:alnum:]]*`

Floating points values (referred to as DVALUE) must belong to

`[+|-]? [:digit:]+ [\. [:digit:]+ [[e|E] [+|-]? [:digit:]+]?]?`

Please refer to the flex sources in `3D-ICE/flex` for more details. It is also possible to insert comments at the end of a line (`//`) or to comment an entire block (`/* ... */`) of the input files (similar to C or C++).

6. Creating a 3D-ICE project

As mentioned in Section 5.B, a 3D-ICE project consists of writing a stack descriptor file and one or more floorplan files. See the example Stack Description File and Floorplan files provided in the `./bin` folder for reference.

A. Stack Description File

The stack description file (*.stk) is a netlist that specifies all the physical and geometrical properties of the 3D-IC for the simulation. The extension of the file is not relevant-it will be parsed independent of its presence or content.

The stack description file contains EIGHT main sections (mandatory and optional) and they must be declared in this order:

1. Materials
2. Heat sink
3. Liquid-cooled cavity
4. Layers
5. Dies
6. Dimensions
7. Stack
8. Analysis options
9. Output instructions

The following description of each of these sections is NOT in the above-mentioned order for ease of presentation and coherence.

i. Materials

The first section of the file contains the list of materials and their properties to be used in the simulation. At least one material must be declared. Materials are declared with the syntax,

```
material MATERIAL_ID :  
  
    thermal conductivity      DVALUE ;  
  
    volumetric heat capacity DVALUE ;
```

where

- `MATERIAL_ID` is a unique identifier to refer to this material,
- `thermal conductivity` is expressed in $\text{W}/\mu\text{m K}$,
- `volumetric heat capacity` is expressed in $\text{J}/\mu\text{m}^3\text{K}$.

Materials declared here but not used in the following sections (channel, dies or stack) will be reported with a warning message (`stderr`).

Example

```
material SILICON :
    thermal conductivity    1.300e-04 ;
    volumetric heat capacity 1.628e-12 ;
```

ii. Layers

A layer is a horizontal section of the stack with a given thickness and made by a given material. A layer declared in this section can be further cited/used when declaring the sequence of stacked elements in the 3D IC later in the file. Layers declared in this optional section of the file can be referred only in the stack section but not when declaring a die. Layers can be declared with the following syntax:

```
layer LAYER_ID :
    height      DVALUE ;
    material     MATERIAL_ID ;
```

where

- LAYER_ID is a unique identifier to refer to this layer,
- height is expressed in μm ,
- MATERIAL_ID is the (previously declared) identifier of the material composing the layer.

Example

```
layer PCB :
    height 10 ;
    material BEOL ;
```

iii. Dies

A die is a group of layers stacked together to form a single entity that is used when declaring the sequence of stacked elements of the 3D IC later in the file. This can represent an actual IC die in the stack. You can declare multiple dies, and use a single die multiple times during the stack description. The Dies section is mandatory and must contain at least one die element. A die must contain one *source* layer (the term *source* layer is used to denote those layers of the stack which contain active electronic components, and hence, provide the heat source for the simulation) and zero or more passive layers. The source layer can be placed at any location in the stack of layers in a die.


```
die DIE_ID :
    [layer IVALUE MATERIAL_ID ; ]*
    source IVALUE MATERIAL_ID ;
    [layer IVALUE MATERIAL_ID ; ]*
```

where

- DIE_ID is the unique identifier used to refer to the declared die,

- IVALUE is the height of the layer (in μm),
- MATERIAL_ID is the (previously declared) identifier of the material composing the layer.

The order of the layers within the die reflects their vertical disposition in the 3D IC, i.e., the first layer declared is the top most layer in the die (closer to the ambient) while the last one is the one at the bottom (closer to the PCB). Two examples of die declaration and their illustrations (not to scale) are shown below.



iv. Heat Sink

This is an optional section which includes the air-cooled heat sink in the 3D IC. All the faces of the 3D IC stack are modeled as adiabatic walls by default. When the Heat Sink is specified, the top surface of the stack is connected to the ambient via a thermal resistance. Two different types of heat sink can be declared:

connection to ambient :

heat transfer coefficient DVALUE ;

ambient temperature DVALUE ;

or

heat sink :

sink height DVALUE, area DVALUE, material MATERIAL_ID ;

spreader height DVALUE, area DVALUE, material MATERIAL_ID ;

heat transfer coefficient DVALUE ;

ambient temperature DVALUE ;

where

- `sink height` and `spreader height` are the heights of the corresponding layers, in μm .
- `sinkarea` and `spreader area` are the horizontal surface of the corresponding layers, in μm^2 .
- `MATERIAL_ID` is the identifier of the material that composes the corresponding layer.
- `heat transfer coefficient` of the heat sink is expressed in $\text{W}/\mu\text{m}^2\text{K}$,
- `ambient temperature` is the ambient temperature expressed in K.

The first type of heat sink, declared as “`connection to ambient`” connects the nodes of the layer at the top of the 3D stack directly to the environment.

Example

```
connection to ambient:
  heat transfer coefficient 1.0e-7;
  ambient temperature 300;
```




Figure 9: Connection to ambient

The second type of heat sink, declared as “`heat sink`”, gives the option to add to the 3D stack a spreader layer and a heat sink. For both the objects, thickness (height), area and material must be specified. In this case, the resistances that connect the top most layer of the stack to the environment are computed following the method described in [7].

Example

```
heat sink:
  sink    height 2e3 ,
          area    16e6 ,
  material COPPER ;
  spread height 1e3 ,
          area    2.25e6 ,
  material CERAMIC ;
  heat transfer coefficient 1.0e-7;
  ambient temperature 300;
```




Figure 10: Heat sink

v. Liquid-cooled cavity

The Liquid-cooled cavity(henceforth referred to simply as “cavity”) section provides 3D-ICE information about the interlayer liquid-cooled heat sink. Since in a given 3D IC design, all interlayer cavities are designed identically, only one declaration of (and one type of) cavity is allowed. As mentioned earlier, there are four types of cavities in 3D-ICE and the flow of coolant is always the SOUTH-NORTH or +y direction. This section can be omitted for simulating a 3D IC without liquid cooling (solid only).

Microchannel 4-resistor model

```
microchannel4rm :  
  
    height DVALUE ;  
  
    channellengthDVALUE ;  
  
    wall lengthDVALUE ;  
  
    [first wall length DVALUE ; ]?  
  
    [last wall length DVALUE ; ]?  
  
    wall material MATERIAL_ID ;  
  
    coolant flow rate DVALUE ;  
  
    coolant heat transfer coefficient [ DVALUE | side DVALUE ,  
                                     top DVALUE ,  
                                     bottom DVALUE ] ;  
  
    coolant volumetric heat capacity DVALUE ;  
  
    coolant incoming temperature DVALUE ;
```

where

- `height` (in μm) corresponds to the height of the microchannel layer. This must exactly correspond to the microchannel height from the surface of the top wall to the surface of the bottom wall because of the 4RM-CTTM and 2RM-CTTM modeling requirements. Solid walls bounding the top and bottom faces of the microchannel would constitute new layers that should be declared separately,
- `channellength` and `walllength` are the cross sectional lengths of the channel and the wall (in the x direction, all in μm). During the discretization of the system (see the Dimensions section for more details), 3D-ICE automatically starts with a wall at the eastern most end of the layer, and alternate channels and walls along the x direction. But the user must ensure that dimensions are such that the layer always ends with a wall (in other words, the number of columns must always be an odd number),
- `first wall length` and `last wall length` (in μm) are optional properties that represents the length of the western-most (the first column) and eastern-most (the last column) walls. This option has been included since, during the fabrication of microchannels on the back of the substrate, although the etching mask pattern is predominantly regular, there are chances of irregularities at the ends, or there might be a deliberate use of different dimensions for the first and the last wall to preserve uniformity and symmetry of heat transfer coefficient. If one of these two dimensions (or both) is not declared, then the `wall length` will be used at its corresponding location,
- `MATERIAL_ID` is the identifier of the material composing the walls (the ID must be previously declared in the materials section),

- coolant flow rate is expressed (in ml/min) and it refers to the volume of coolant flowing per unit time per channel layer (cavity) in the stack. If you have multiple layers of microchannels, the total flow rate must be divided by the number of channel layers and given as a single input,
- coolant heat transfer coefficient is the Heat Transfer Coefficient of convective heat removal from the walls into the coolant (in $W/\mu m^2 K$). It is possible to specify a single value of HTC for all the wetted surfaces of the microchannel or specify three different values- one for the *side* wall surfaces, one for the *top* and one for the *bottom* wall surface of the microchannel.
- coolant volumetric heat capacity is expressed in $J/\mu m^3 K$,
- coolant incoming temperature is the inlet coolant temperature expressed in K.

Example

```

microchannel4rm:
  height 100 ;
  channel length 50;
  wall length 50;
  first wall length 25;
  last wall length 25;
  wall material SILICON;
  coolant flow rate 42;
  coolant heat transfer coefficient
    side 2.7132e-8,
    top 5.7132e-8,
    bottom 4.7132e-8;
  coolant volumetric heat capacity
    4.172e-12;
  coolant incoming temperature
    300;

```




Figure 11: Microchannel cavity layer

Microchannel 2-resistor model

```

microchannel2rm :

  height DVALUE ;

  channel length DVALUE ;

  wall length DVALUE ;

  wall material MATERIAL_ID ;

  coolant flow rate DVALUE ;

  coolant heat transfer coefficient [ DVALUE | top DVALUE ,
                                     bottom DVALUE ] ;

  coolant volumetric heat capacity DVALUE ;

  coolant incoming temperature DVALUE ;

```

Note that the only change from microchannel-4rm model is the omission of the first- and the last-wall widths and the side-way heat transfer coefficient, as was described in Section 5.A.ii. Also, remember that these heat transfer coefficients are effective “projections” of the heat transfer coefficients on the top and the bottom surfaces of the cavity. If you have the top, bottom and side heat transfer

coefficients (HTCs) for a microchannel structure, you can compute these effective HTCs using the following formulae [2]:

$$HTC_{top,eff} = \frac{HTC_{top} \times \text{channel width} + HTC_{side} \times \text{channel height}}{\text{channel pitch}}$$

$$HTC_{bottom,eff} = \frac{HTC_{bottom} \times \text{channel width} + HTC_{side} \times \text{channel height}}{\text{channel pitch}}$$

Example

```
microchannel2rm:
  height 100 ;
  channel length 50;
  wall length 50;
  wall material SILICON;
  coolant flow rate 42;
  coolant heat transfer coefficient
    top 5.5698e-8,
    bottom 5.0698e-8;
  coolant volumetric heat capacity
    4.172e-12;
  coolant incoming temperature
    300;
```




Figure 12: Microchannel cavity layer

Pinfins

```
pinfin:
  height DVALUE ;
  pin diameter DVALUE ;
  pin pitch DVALUE ;
  pin distribution [inline | staggered] ;
  pin material MATERIAL_ID ;
  darcy velocity DAVLUE ;
  coolant volumetric heat capacity DVALUE;
  coolant incoming temperature DVALUE;
```

where

- `height` (in μm) corresponds to the height of the pinfin layer. This must exactly correspond to the height of the pins in accordance with the 2RM-CTTM model,
- `pin diameter` and `pin pitch` are diameter and the pitch (all in μm) of the pins. 3D-ICE uses this information to compute the pin density and the porosity of the model,
- `pin distribution` defines the type of pinfin distribution- `inline` for inline (Fig. 4) and `staggered` for staggered (Fig. 5),
- `MATERIAL_ID` is the identifier of the material composing the pins (the ID must be previously declared in the materials section),

- darcy velocity is expressed (in $\mu\text{m}/\text{sec}$) and it refers to the darcy velocity in the cavity. This is used instead of the coolant flow rate because darcy velocities are more commonly measured quantities in experiments and CFD simulations due to their geometry independence, and more useful while expressing the heat transfer coefficient correlations in empirical studies.

Example (Pinfins inline)

```
pinfin:
  height 100 ;
  pin diameter 50 ;
  pin pitch 100 ;
  pin distribution inline ;
  pin material silicon ;
  darcy velocity 1.1066e+06 ;
  coolant volumetric heat capacity
  4.172638e-12 ;
  coolant incoming temperature
  300.0 ;
```




Figure 13: Pinfin inline cavity layer

Example (Pinfins staggered)

```
pinfin:
  height 100 ;
  pin diameter 50 ;
  pin pitch 100 ;
  pin distribution staggered ;
  pin material silicon ;
  darcy velocity 1.1066e+06 ;
  coolant volumetric heat capacity
  4.172638e-12 ;
  coolant incoming temperature
  300.0 ;
```




Figure 14: Channel layer

Note: If neither the conventional nor the liquid-cooled heat sink is declared in the Stack DescriptionFile, then the parsing will end with a warning message. Note that this would cause the temperatures to blow up unbounded with time in the presence of non-zero heat sources.

vi. Stack

This section builds the vertical structure of the stack. The stack is composed of Dies (as previously declared) and layers and/or channels.

```
stack :

  [ layer  LL_ID LAYER_ID ;          |
    channel CC_ID ;                  |
    die    DD_ID  DIE_ID floorplan "PATH" ; ] +
```

where

- LL_ID, CC_ID and DD_ID are identifiers used to name the stack elements and they can be used in the simulator code to refer to the corresponding element. They must be unique for each element,

- `LAYER_ID` is the identifier of a layer (as previously declared).
- `DIE_ID` is the identifier of a die (as previously declared) and `PATH` is the path to the floorplan file. This floorplan will be placed on the declared source layer in the definition of the die. The floorplan files contain information of the location and power dissipation activity of various floorplan components for the given die (see the description of Floorplan Files for more details). The same `DIE_ID` can be used multiple times (with different identifiers `DD_ID`) in a stack with the same or different floorplans, if identical/similar dies exist in a single IC.

Note that the keyword “channel” is used to specify the cavity in the stack- irrespective of which type of interlayer cavity (microchannels or pin fins) was declared in the stack descriptor.

The above grammar for the stack section is not the representative of the real syntax that is supported and is only given for simplicity. For instance, the three main elements (layer, die and channel) that compose the stack can be declared in any order but the final stack sequence must satisfy the following constraints:

- there must be at least one die
- it cannot begin or end with a cavity (i.e., liquid-cooled cavities can’t be the bottommost or the topmost layers in a stack)
- there cannot be two consecutive cavities
- cavities can be used only if previously declared
- layers are optional

Declaring a layer in a stack is not mandatory but we left this option to support stacks with irregular patterns of dies and channels or to build auxiliary layers (such as a bonding layer).

As in the case of defining dies, the final sequence of stack elements reflects their vertical disposition. The first stack element corresponds to the topmost element while the last element declared is closest to the PCB.

Example

```
material SILICON :
    thermal conductivity    1.30e-4;
    volumetric heat capacity 1.628e-12;
material BEOL :
    thermal conductivity    2.25e-6;
    volumetric heat capacity 2.175e-12;

connection to ambient:
    heat transfer coefficient 1e-07 ;
    ambient temperature 300 ;

microchannel4rm:
    height 100 ;
    channel length 50;
    wall length50;
    first wall length 25;
    last wall length 25;
    wall material SILICON;
    coolant flow rate 42;
    coolant heat transfer coefficient
        top 5.7132e-8,
        bottom 4.7132e-8;
    coolant volumetric heat capacity
        4.172e-12;
    coolant incoming temperature 300;

layer PCB:
    height 10 ;
    material BEOL ;

die TOP_IC:
    source 2 SILICON;
    layer 50 SILICON;

die BOTTOM_IC:
    layer 10 BEOL;
    source 2 SILICON;
    layer 50 SILICON;

dimension :
    chip length 10000, width 10000;
    cell length 50, width 50;

stack:
    die      MEMORY_DIE      TOP_IC      floorplan "./mem.flp";
    channel TOP_CHANNEL;
    die      CORE_DIE        BOTTOM_IC   floorplan "./core.flp";
    channel BOTTOM_CHANNEL;
    layer    BOTTOM MOST     PCB ;
```




Figure 15: Complete stack

vii. Dimensions

This section of the Stack Description File declares the xy dimensions of the entire chip and the discretization sizes for the thermal cells (all in μm).

```
dimension :

    chip length DVALUE , width DVALUE ;

    cell length DVALUE , width DVALUE ;
```

The entire chip is discretized based on the same cell length (along x direction) and width (along y direction) values. The discretization along the z direction is not specified, since the height of a thermal cell is taken to be the same as the height of the layer in which it exists as shown for a layer in Fig. 16. However, if you want a finer discretization than that along the z direction, then you will have to split the layer into multiple layers of the same material stacked on the top of each other in the declaration of die/stack, as shown in Fig. 17 (here $h_1+h_2=h$).




Figure 16: Discretization of a single layer




Figure 17: Discretization of a single layer split into 2 layers

Example

```
dimension :
  chip length 10000, width 10000;
  cell length 50, width 50;
```

If the microchannel 4-resistor model is used, the cell length is determined solely based on the cross sectional dimensions of channel and wall in the Channel section of the Stack Description File. This is because the 4RM-CTTM modeling used in 3D-ICE requires that the entire cross section of the microchannel be a part of the thermal cell. This means that the cell length at some position along the x direction in the model is dependent upon the channel or the wall cross sectional width at that point. Note that given different values of `channel length`, `wall length`, `first wall length` and `last wall length`, the discretization along this direction will be non-uniform in the simulator. For all other liquid-cooled cavity types, the cell length and cell width are necessary and sufficient to create a uniform discretization of the entire structure.

For the purpose of illustration, the discretized domain corresponding to the stacked structure built in Fig. 15 is shown in Fig. 18. Note that the `cell length` must still be declared in all cases, in spite of the fact that it will be ignored during the parsing when microchannel 4-resistor model cavity is used.




Figure 18: Discretized computational domain for the stack shown in Fig. 15

Problem Complexity

Remember that the dimensions of the chip, together with the dimensions of the thermal cells will directly affect the performance of the simulation. Indeed, the number of cells (nodes) influences both the amount of memory used by the library and the time needed to solve the linear system[1]. The main computational effort of the simulator is incurred during the execution of SuperLU and `blas` libraries. Specifically, the LU factorization of the system matrix is the most time/memory intensive. Hence, for large problem sizes, the availability of memory must be ensured to prevent this step from failing during the simulation.

viii. Analysis options

In 3D-ICE either transient or steady state simulations can be performed.

```
solver:
    [steady | transient step DVALUE, slot DVALUE] ;
    initial temperature DVALUE ;
```

where

- `steady` and `transient` indicate steady state and transient analysis respectively,
- `step` corresponds to the internal stepping time (in sec) to be used in the transient simulation in 3D-ICE,
- `slot` is the slot time (in sec) for which each power value in the floorplan file lasts for in the transient simulation (see Section 6.B). This value must be greater or equal to the stepping time value,
- `initial temperature` denotes the initial temperature (in K) for the simulation.

Example (steady state)

```

solver :
    steady;
    initial temperature 300.0;

```

Example (transient)

```

solver :
    transient step 0.02, slot 0.2;
    initial temperature 300.0;

```

ix. Output Instructions

A variety of outputs can be obtained from 3D-ICE simulations, a temperature and power. In this last section of the stack descriptor, various output instructions can be provided for this purpose. The syntaxes for these instructions follow the model of conventional circuit simulators.

output:

```

[ T( VER_ID, DVALUE, DVALUE, "PATH" [, INSTANCE_ID]? ) ; |
Tflp( DD_ID, "PATH", OUTTYPE_ID [, INSTANCE_ID]? ) ; |
Tflpel( DD_ID.FLPEL_ID, "PATH", OUTTYPE_ID [, INSTANCE_ID]? ) ; |
Tmap( VER_ID, "PATH" [, INSTANCE_ID]? ) ; |
Pmap(DD_ID, "PATH" [, INSTANCE_ID]? ) ; |
Tcoolant( CC_ID, "PATH", OUTTYPE_ID [, INSTANCE_ID]? ) ; ]+

```

where

- **T** is an instruction to print the temperature of a particular thermal cell, identified by its three-dimensional coordinates in the thermal grid, in a text file. The output format consists of two columns- the time instance (in sec) against the corresponding temperature value (in K).

T(VER_ID, DVALUE, DVALUE, "PATH" [, INSTANCE_ID]?) ;

- VER_ID specifies the “z” location of the thermal cell. This can assume a value of **LL_ID**, **DD_ID** or a **CC_ID** (see Section 6.A.v). If a die (**DD_ID**) is specified, then the thermal cell would be located in the source layer of that die,
- D_VALUE, D_VALUE (in μm) specify the “x” and “y” locations of the thermal cell. An error message is printed if the location specified is outside of the computational domain,
- PATH is the path of the text file in which the output must be written,
- INSTANCE_ID is an optional parameter that can assume a value of **step**, **slot** or **final** and specifies the frequency in which the output must be reported- at the end of each internal time step, at the end of each time slot, or only at the end of the simulation respectively. Note that in the case of steady state simulation, only those instructions that have the option **final** are executed. When no parameter is specified, **final** is assumed,

- `Tflp` is an instruction to print the temperature of a particular floorplan (the source layer of a particular die), identified by its die identifier (`DD_ID`). The output format consists of two columns- the time instance (in sec) against the corresponding temperature value (in K).

```
Tflp( DD_ID, "PATH", OUTTYPE_ID [, INSTANCE_ID]? ) ;
```

- `DD_ID` is the identifier of the die whose floorplan temperature must be printed.
 - `PATH` is the path of the text file in which the output must be written,
 - `OUTTYPE_ID` specifies the exact nature of the temperature to be reported, and can assume the values `maximum`, `minimum` or `average`, corresponding maximum, minimum and average temperature in the floorplan respectively,
 - `INSTANCE_ID` is an optional parameter that can assume a value of `step`, `slot` or `final` and specifies the frequency in which the output must be reported- at the end of each internal time step, at the end of each time slot, or only at the end of the simulation respectively. Note that in the case of steady state simulation, only those instructions that have the option `final` are executed. When no parameter is specified, `final` is assumed.
- `Tflpel` is an instruction to print the temperature of a particular floorplan element (the power block of the source layer of a particular die), uniquely identified by its die identifier (`DD_ID`) and floorplan element identifier (`FLPEL_ID`). The output format consists of two columns- the time instance (in sec) against the corresponding temperature value (in K).

```
Tflpel( DD_ID.FLPEL_ID, "PATH", OUTTYPE_ID [, INSTANCE_ID]? ) ;
```

- `DD_ID.FLPEL_ID` contains the die and the floorplan element identifier (see Section 6.B) which can uniquely identify a floorplan element in the entire stack,
 - `PATH` is the path of the text file in which the output must be written,
 - `OUTTYPE_ID` specifies the exact nature of the temperature to be reported, and can assume the values `maximum`, `minimum` or `average`, corresponding maximum, minimum and average temperature in the floorplan element area respectively,
 - `INSTANCE_ID` is an optional parameter that can assume a value of `step`, `slot` or `final` and specifies the frequency in which the output must be reported- at the end of each internal time step, at the end of each time slot, or only at the end of the simulation respectively. Note that in the case of steady state simulation, only those instructions that have the option `final` are executed. When no parameter is specified, `final` is assumed.
- `Tmap` is an instruction to print the temperature map of a particular layer in the stack, identified by its identifier. The output is printed in a matrix format in a text file, with each line representing a row of thermal cells in the thermal grid (counted in “y” direction), and every temperature value (in K) in each line corresponding to a thermal cell (columns, counted in the “x” direction). Once the temperature map at a given time point is printed in this format, the temperature map at the next time point requested is printed from the next line in the same format. Hence, the

total number of lines in the file is the product of the number of rows in the thermal grid and the number of instances the output is printed. In addition, whenever the instruction `Tmap` is used in a project, two other files, named `"x_axis.txt"` and `"y_axis.txt"` are printed which consist of the indices of the columns and rows (in μm), respectively, for these temperature maps. They can be used to visualize the data in the text files with the help of applications supporting graphical outputs such as Matlab.

```
Tmap( VER_ID, "PATH" [, INSTANCE_ID]? ) ;
```

- `VER_ID` identifies the layer for which the temperature map must be printed. This can assume a value of `LL_ID`, `DD_ID` or a `CC_ID` (see Section 6.A.v). If a die (`DD_ID`) is specified, then the source layer of that die is used,
 - `PATH` is the path of the text file in which the output must be written,
 - `INSTANCE_ID` is an optional parameter that can assume a value of `step`, `slot` or `final` and specifies the frequency in which the output must be reported- at the end of each internal time step, at the end of each time slot, or only at the end of the simulation respectively. Note that in the case of steady state simulation, only those instructions that have the option `final` are executed. When no parameter is specified, `final` is assumed.
- `Pmap` is an instruction to print the power map of a particular die in the stack, identified by its identifier. The output is printed with the same criteria implemented to print a temperature map.

```
Pmap(DD_ID, "PATH" [, INSTANCE_ID]? ) ;
```

- `DD_ID` identifies the die for which the power map must be printed.
 - `PATH` is the path of the text file in which the output must be written,
 - `INSTANCE_ID` is an optional parameter that can assume a value of `step`, `slot` or `final` and specifies the frequency in which the output must be reported- at the end of each internal time step, at the end of each time slot, or only at the end of the simulation respectively. Note that in the case of steady state simulation, only those instructions that have the option `final` are executed. When no parameter is specified, `final` is assumed.
- `Tcoolant` is an instruction to print the temperature of the coolant at the outlet of a cavity identified by its channel identifier (`CC_ID`). The output format consists of two columns- the time instance (in sec) against the corresponding temperature value (in K).

```
Tcoolant( CC_ID, "PATH", OUTTYPE_ID [, INTANCE_ID]? ) ;
```

- `CC_ID` is the identifier of the cavity whose outlet temperature must be printed,
- `PATH` is the path of the text file in which the output must be written,
- `OUTTYPE_ID` specifies the exact nature of the temperature to be reported, and can assume the values `maximum`, `minimum` or `average`, corresponding maximum, minimum and average outlet temperature respectively,
- `INSTANCE_ID` is an optional parameter that can assume a value of `step`, `slot` or `final` and specifies the frequency in which the output must be reported- at the end of each internal time step, at the end of each time slot, or only at the end of the simulation

respectively. Note that in the case of steady state simulation, only those instructions that have the option `final` are executed. When no parameter is specified, `final` is assumed.

Examples

Two example stack descriptor files, one steady state and one transient, based on the above discussions are shown below.

Example (steady state)

```
material SILICON :
    thermal conductivity    1.30e-4;
    volumetric heat capacity 1.628e-12;
material BEOL :
    thermal conductivity    2.25e-6;
    volumetric heat capacity 2.175e-12;

conventional heat sink:
    heat transfer coefficient 1.0e-7;
    ambient temperature 300;

microchannel4rm:
    height 100 ;
    channel length 50;
    wall length50;
    first wall length 25;
    last wall length 25;
    wall material SILICON;
    coolant flow rate 42;
    coolant heat transfer coefficient top 5.7132e-8,
                                   bottom 4.7132e-8;
    coolant volumetric heat capacity4.172e-12;
    coolant incoming temperature300;
layer PCB:
    height 10 ;
    material BEOL ;
die TOP_IC:
    source 2 SILICON;
    layer 50 SILICON;
die BOTTOM_IC:
    layer 10 BEOL;
    source 2 SILICON;
    layer 50 SILICON;
dimension :
    chip length 10000, width 10000;
    cell length 50, width 50;

stack:
    die      MEMORY_DIE      TOP_IC      floorplan"./mem.flp";
    channel TOP_CHANNEL;
    die      CORE_DIE        BOTTOM_IC floorplan"./core.flp";
    channel BOTTOM_CHANNEL;
    layer    BOTTOM_MOST      PCB ;
```

```

solver:
    steady ;
    initial temperature 300.0 ;

output:
    T( MEMORY_DIE, 5000, 3000, "output1.txt",final);
    Tmap( CORE_DIE, "output2.txt", final );
    Pmap( CORE_DIE, "output3.txt", final );
    Tflp1(CORE_DIE.core1, "output4.txt", average);
    Tcoolant(TOP_CHANNEL, "output5.txt", maximum);

```

Example (transient)

```

material SILICON :
    thermal conductivity    1.30e-4;
    volumetric heat capacity 1.628e-12;
material BEOL :
    thermal conductivity    2.25e-6;
    volumetric heat capacity 2.175e-12;

conventional heat sink:
    heat transfer coefficient 1.0e-7;
    ambient temperature 300;

microchannel4rm:
    height 100 ;
    channel length 50;
    wall length 50;
    first wall length 25;
    last wall length 25;
    wall material SILICON;
    coolant flow rate 42;
    coolant heat transfer coefficient top 5.7132e-8,
        bottom 4.7132e-8;
    coolant volumetric heat capacity 4.172e-12;
    coolant incoming temperature 300;

layer PCB:
    height 10 ;
    material BEOL ;
die TOP_IC:
    source 2 SILICON;
    layer 50 SILICON;
die BOTTOM_IC:
    layer 10 BEOL;
    source 2 SILICON;
    layer 50 SILICON;

dimension :
    chip length 10000, width 10000;
    cell length 50, width 50;

stack:
    die      MEMORY_DIE      TOP_IC      floorplan `./mem.flp`;
    channel TOP_CHANNEL;
    die      CORE_DIE        BOTTOM_IC floorplan `./core.flp`;
    channel BOTTOM_CHANNEL;
    layer    BOTTOM_MOST      PCB ;

```

```

solver:
    transient step 0.02, slot 0.2 ;
    initial temperature 300.0 ;

output:
    T( MEMORY_DIE, 5000, 3000, "output1.txt",step);
    Tmap( CORE_DIE, "output2.txt", slot );
    Pmap( CORE_DIE, "output3.txt", slot );
    Tflp( MEMORY_DIE, "output4.txt", minimum, step);
    Tflpel(CORE_DIE.core1, "output5.txt", average, final);
    Tcoolant(TOP_CHANNEL, "output6.txt", maximum);

```

B. Floorplan File

Every die in the stack must be related to a "Floorplan File" (*.flp), which essentially provides the power dissipation profile (or heat sources) for the simulation. Each Floorplan file must contain the list of functional blocks (cores, caches, memories, etc), their positions, and the power dissipation as a function of time.

Every functional block, here called *floorplanelement*, is an area inside the die, laid out in the source layer. Each floorplan element has a unique identifier- the name it is assigned. In addition, the position and the dimensions of each floorplan element are given (in μm) based on the same Cartesian coordinates that was used for building the stack, with the origin at the SOUTH-WEST corner of the source layer. An example floorplan of a 1cmX1cm die with the reference coordinates is shown in Fig. 19. All the distances shown here are in μm .

A floorplan element in the Floorplan File is declared using the following syntax.

```

IDENTIFIER :

    position  DVALUE , DVALUE ;

    dimension DVALUE , DVALUE ;

    [ power values DVALUE [ , DVALUE ]* ]? ;

```

Or, if its surface cannot be represented as a simple rectangle, the following syntax is also supported:

```

IDENTIFIER :

    [ rectangle (DVALUE , DVALUE , DVALUE , DVALUE ) ; ]+

    [ power values DVALUE [ , DVALUE ]* ]? ;

```

where

- IDENTIFIER is the unique identifier used to name the floorplan element. This string must be unique within the floorplan file it belongs to but it can be used on a different file,
- position, is the (x,y) coordinate of the SOUTH-WEST corner of the floorplan element (in μm),
- dimension is the (length, width) dimensions of the floorplan element (in μm),

- `rectangle` is the (x,y, length, width) description of a part of the floorplan element (in μm)
- The `DVALUE(s)` against the keyword `power values` are the list of power dissipation values (expressed in W) of the floorplan element for each *time slot* (scroll down for the explanation of time slots in 3D-ICE) separated by commas.

As an example, the declaration for the Core 0 and the L2 Chache 0 in Fig. 19 with 5 time slots would be as follows:

Example

```
L2_Cache_0 :
    position 0, 1750;
    dimension 1250, 7500;
    power values 0.3, 0.3, 0.4, 0.2, 0.3;

Core_0 :
    rectangle ( 0, 0, 1250, 1750) ;
    rectangle (1250, 0, 3750, 2000) ;
    rectangle (1250, 2000, 2500, 1500) ;
    power values 0.5, 0.5, 0.1, 0.1, 0.7;
```




Figure 19: An example floorplan for a 1cmX1cm IC die. Notice the irregular floorplans Core 0, Core 1, Core 2, Core 3.

The following points must be kept in mind while writing the floorplan files:

- Floorplan elements must not overlap. During the parsing of the Floorplan File, the values describing the elements are checked to verify that all the elements are inside the chip and that they do not overlap.
- When the stack structure is discretized based on the given thermal cell dimensions, the power dissipated by a floorplan element is **uniformly** divided among the thermal cells that *intersect its surface*. It is **not required** that the position and dimensions of the floorplan elements are an

exact multiple of the dimensions of the thermal cells. Therefore, depending on the layout of the IC, a thermal cell can receive the power from different floorplan elements. For the same reason, a floorplan element can also be smaller than the thermal cell itself. The assignment of thermal cells to different floorplan elements is illustrated in Fig. 20. The thermal cell highlighted in the center of the image receives as input the power coming from the FPU unit, the L1 Cache 1 and the Core 1 at the same time. The power value consumed by each component will be scaled according to the fraction of the surface of the thermal cell covered by the corresponding IC unit:

$$Input = Power_{FPU} \frac{a_1}{Area_{FPU}} + Power_{CacheL1} \frac{a_2}{Area_{CacheL1}} + Power_{Core1} \frac{a_3}{Area_{Core1}}$$

- The same Floorplan File can be assigned to 2 or more dies in the Stack Description File if they happen to have identical structure and behavior in the design. Each floorplan element in the entire design, in that case, is uniquely identified by the `DD_ID` identifier of the corresponding die element and the `IDENTIFIER` of the floorplan element.
- If two dies in the stack have the same floorplan but during the simulation they have a different power dissipation activity, then 2 different Floorplan Files must be created for each die and assigned to the corresponding die element declarations in the Stack Description File. This is because the power dissipations are directly linked to each floorplan element in the Floorplan File.
- It is possible to have gaps in the floorplan- regions where there is no floorplan element and hence, no power dissipation. The thermal cells in these regions will simply not be assigned any source value during the solving of system equations.




Figure 20: Assignment of thermal cells to floorplan elements

Time Slots

The entire time-interval of simulation (ToS) in 3D-ICE is divided into *time slots*- the minimum time duration for which the switching activity of the floorplan elements has been resolved. For example, if for a given design the switching activity (a measure of how much a floorplan component is active, directly

related to its power dissipation) is sampled every 200ms during a 1 second ToS, then there are 5 time slots for the 3D-ICE simulation. And hence, there must be 5 values of power dissipation for each floorplan element declaration in the Floorplan File. Conversely, the number of power values for the floorplan element is interpreted as the number of time slots (NoTS) for the simulation. In all 3D-ICE simulations, it is assumed that the power dissipation in a particular thermal cell is CONSTANT during the period of a time slot- calculated based on the power value of the corresponding floorplan element, in which the thermal cell exists, for that time slot (see Fig. 21).

Note that time slot is different from *time step*, which is the discretization length of the time-domain for the numerical integration of the system of differential equations representing the entire thermal circuit created inside 3D-ICE. The exact durations of both time slot and time step are given when running the simulator. For further details, see Chapter 5.




Figure 21: Power dissipation profile of Core 0 and Core 6 in Fig. 19

Since uniform sampling of the switching activity for all the components in an IC is assumed, the number of power values given for each floorplan element in the Floorplan File should be the same. Even when there is zero power dissipation for a particular floorplan element during some time slot (Core 0 at slot #4 in Fig. 21) or if a particular floorplan element has constant power dissipation during some or all time slots (Core 6 during slot #2 and slot #3 in Fig. 21), the corresponding power values must be mentioned explicitly for each time slot. The function `emulate_slot` returns a value indicating the end of the simulation whenever it finds a floorplan element with an empty list of power values. Therefore, the ToS will be determined by the length of the shortest list of power values within all the floorplan elements that are declared in the dies.

7. Network interface for remote simulations

The 3D-ICE software library includes software object that can be used to implement simulations where two distinct processes, a client and a server, communicate through a network socket. In this interface, 3D-ICE running on a host machine, acts as the server and the some device or external source that generates the power trace inputs acts as a client. Such set of data structures and functions can be used whenever it is necessary to decouple the generation of power traces and the control of the simulation from its execution, which might requires more computational resources. The network interface can be used in the following scenarios:

- Hardware/software co-simulation: for example, obtaining real-time temperature data from an architecture being designed and prototyped on an FPGA, testing run-time thermal management techniques that depend on data from temperature sensors in a real device. For more details, see [6]
- Creating graphical interfaces that involve communication with other processes running parallel on a computer.
- Integration in other tools or libraries: for example, floorplanning software, software emulators for power traces on the same computer that is running 3D-ICE.

In particular, it will be useful whenever power traces are not known a priori and cannot be written as input for the simulation directly in the floorplan file, and for cases where the power traces are too long and must be read from large data files. The primary project files for a simulation, i.e. the .stk and the flp files, describing the structure of the IC and the floorplan architecture must still be written and stored in the server. The client would simply send the power traces that are otherwise entered in the .flp files. The server performs the thermal simulation and provides the real-time values of the outputs requested in the .stk file. These values will be sent over a socket back to the client, instead of being printed in a text file. The client can also send command to the server to request any kind of change of the state of the simulation (reset temperatures, change flow rate, run simulation steps, etc).

For the usage of the interface, two data types have been added as well as functions to handle them as done for the other components of the library. These data types are `Socket` and `NetworkMessage` and they will be described in Section 9.

8. Running 3D-ICE

In the `3d-ice/bin/` folder, the main simulator file to be compiled is `3D-ICE-Emulator.c`. It has been written to parse and analyze any Stack Description File and the corresponding Floorplan Files placed in the same folder. Once compiled with the make command, the corresponding executable is acts as the thermal simulator application. To simulate a new thermal project, you must:

- create the Stack Description File and the corresponding Floorplan Files according to the instructions in Chapter 6
- run the following command in `3d-ice/bin/`
`$./3D-ICE-EmulatorPATH`

where `PATH` is the path to the Stack Descriptor File containing the description of the 3D-IC project.

The project is then simulated and the results, as requested in the `.stk` file, are written in the text files specified in those instructions by the user.

3D-ICE uses backward Euler method with constant time-stepping to solve the system equations. Hence, the solution is always numerically unconditionally stable. However, accuracy can be increased by reducing the time step value. The local truncation error of backward Euler method behaves as $O(h^2)$, where h is the time step. However, given a ToS, the number of time steps in the entire simulation is a $O(1/h)$ function. Hence, the upper bound of the total accumulated error at the final time point in the simulation behaves as

$$O(h^2) \cdot O(1/h) = O(h).$$

In other words, the total final error is approximately a linear function of the time step. For RC circuits, such as the thermal circuit that 3D-ICE solves, it is common practice to have at least 5 time steps for the duration of a rise time of the output temperature (defined as the time duration for the rise of temperature from 10% to 90% of its steady state value) to resolve the transients accurately.

9. Usage of the 3D-ICE as Software Thermal Library

3D-ICE, in addition to functioning as a stand-alone thermal simulator, can also serve as a software thermal library that can be used to build customized applications. Such requires knowledge about the organization and the use of various functions and data structures that are built into 3D-ICE. With the release of 3D-ICE 2.0, we are happy to announce that a new, useful tool for the visualization of the software library, called Doxygen [5], was used to build an online documentation of the 3D-ICE library. It contains convenient access all information about the functions and the files in the library, including a hierarchical visualization of the various dependencies between them. This new documentation can be accessed by running in the `make doc` command inside the 3D-ICE folder, and then opening `3d-ice/doc/html/index.html` on your web browser. We urge users who are interested in employing 3D-ICE to build custom applications to refer to this documentation.

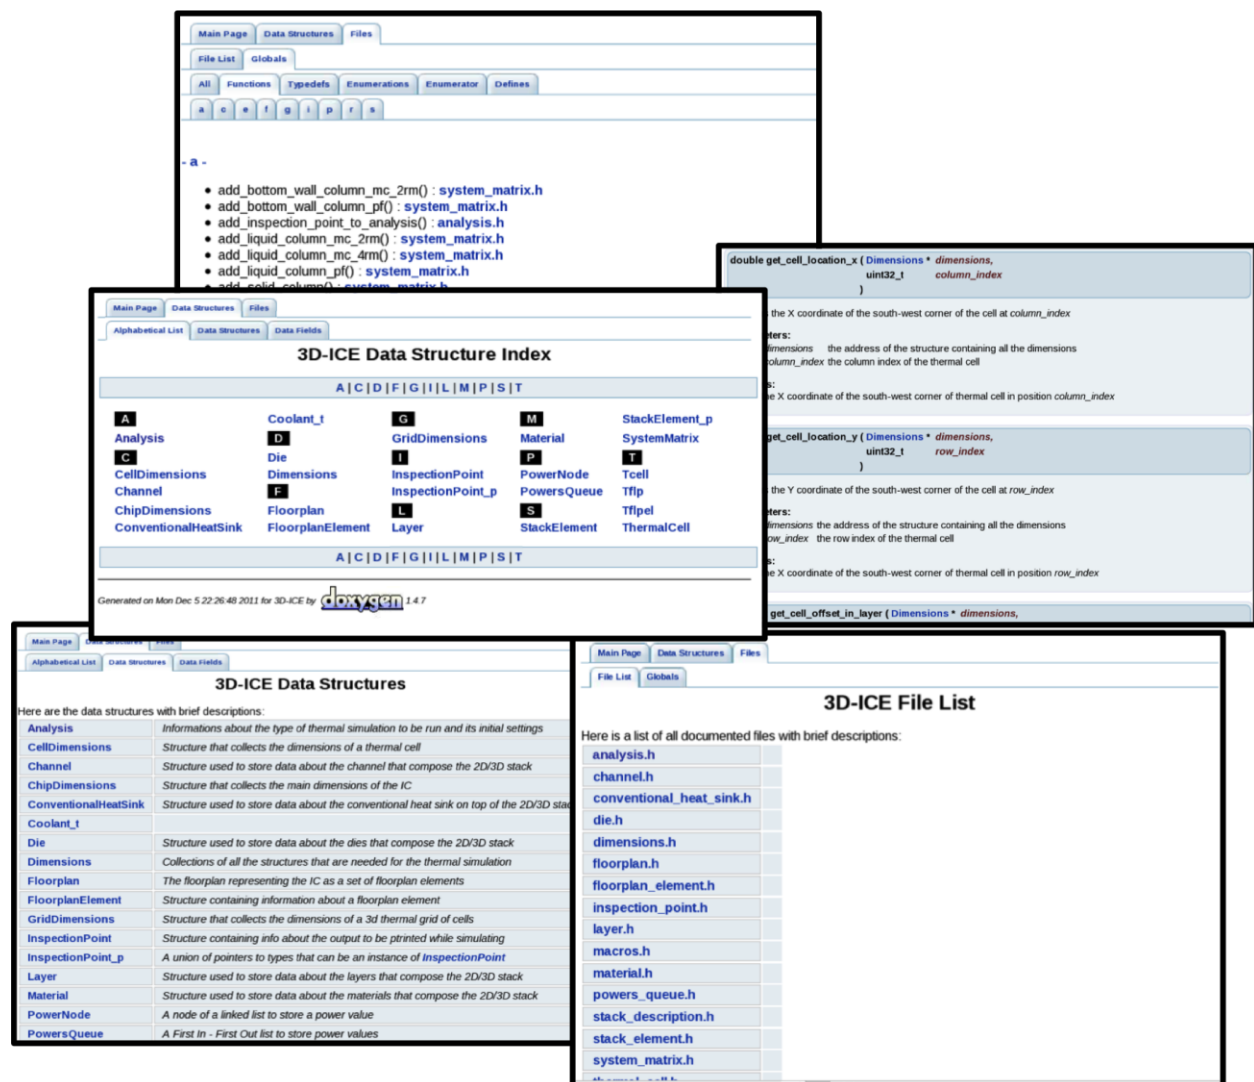


Figure 22: Doxygen documentation for 3D-ICE 2.0

A. StackDescription_t, Analysis_t and Output_t

Using 3D-ICE as a software library involves several steps. First, the problem must be initialized and the necessary data structures containing the 3D IC information must be filled. The three data structures that store the description of the IC stack and the configuration of the simulation to be run are, respectively, `StackDescription_t`, `Analysis_t` and `Output_t`. To use these types, the header files that must be included in the main file to access the homonym data structures are `stack_description.h`, `analysis.h` and `output.h`.

These data structure types collect all the data pertaining to the 3D IC structure and floorplans. An instance of these data types will then be related to the Stack Description File. The functions available to initialize, fill and destroy an instance of these two variables are:

- `stack_description_init (StackDescription_t*)`
- `analysis_init (Analysis_t*)`
- `output_init (Output_t*)`

- `parse_stack_description_file`
 `(String_t, StackDescription_t*, Analysis_t*, Output_t*)`

- `stack_description_destroy (StackDescription_t*)`
- `analysis_destroy (Analysis_t*)`
- `output_destroy (Output_t*)`.

In particular, the function `parse_stack_description_file`, parses the content of a stack file and places the information read into the corresponding data structure. This function can be accessed by including the header file `stack_file_parser.h`. The functions to access the information in this data structure are in the doxygen documentation.

B. ThermalData_t

This data structure type collects all the data needed for the thermal simulation, such as temperatures, power inputs, matrices representing the system equations, etc. Its type is declared in the header `thermal_data.h`. The functions to initialize, build and destroy a `ThermalData_t` variable are:

- `thermal_data_init (ThermalData_t*)`
- `thermal_data_build`
 `(ThermalData_t*, StackElementList_t*, Dimensions_t*, Analysis_t*)`
- `thermal_data_destroy (ThermalData_t*)`.

Just as an instance of `StackDescription_t`, `Analysis_t` and `Output_t` are tied to a Stack Description File, an instance of `ThermalData_t` depends upon the sequence of stack elements in `StackElementList_t` (a member of `StackDescription_t`) and `Analysis_t`. Hence it is necessary to initialize and fill the `StackDescription_t` and `Analysis_t` variables before filling the `ThermalData_t` variable. And for the rest of the simulation project, the three variables must be used in tandem since they refer to the same problem.

C. Emulation and thermal output

Once the data structures mentioned in the previous sections have been filled, the next step is to execute the simulation. It can be done using the functions

- `emulate_step (ThermalData_t*, Dimensions_t*, Analysis_t*)`
- `emulate_slot (ThermalData_t*, Dimensions_t*, Analysis_t*)`
- `emulate_steady (ThermalData_t*, Dimensions_t*, Analysis_t*)`.

`emulate_step` increments the simulation time by a single time step and then terminates- this function can be called iteratively in a loop until the end of the ToS, controlled using the return variables of these functions which indicating whether or not the simulation has reached certain epochs (see the description in the files containing these functions for more details); while `emulate_slot` increments the simulation time until the simulation time of a single power value finishes. Finally, the `emulate_steady` performs the steady simulation (one step simulation).

Outputs can usually be read from the thermal analysis before, during and at the end of the simulation in the text files mentioned in the output instructions written in the .stk file of the project. The structure `Output_t` stores the list of outputs that should be generated.

If you intend to use 3D-ICE as a software library, there are functions in 3D-ICE that enable the extraction of the thermal state of a single floorplan element, an entire floorplan, a single channel outlet or all the coolant outlets in a given channel layer at any time during the simulation. It is also possible to read the temperature of a single thermal cell or print directly the thermal map (a matrix) for a specific layer in the stack. Some of these functions require you to refer to floorplan elements, layers etc. using the corresponding identifiers declared in the Stack Description File. The functions available to generate the output are:

- `generate_output_headers (Output_t *, Dimensions_t *, String_t)`
- `generate_output (Output_t*, Dimensions_t*, Temperature_t*, Source_t*, Time_t, OutputInstant_t)`

The first function creates the text files and must be called only once before the simulation. The function `generate_output`, instead, must be called whenever the output (as the current thermal state of the IC) is needed.

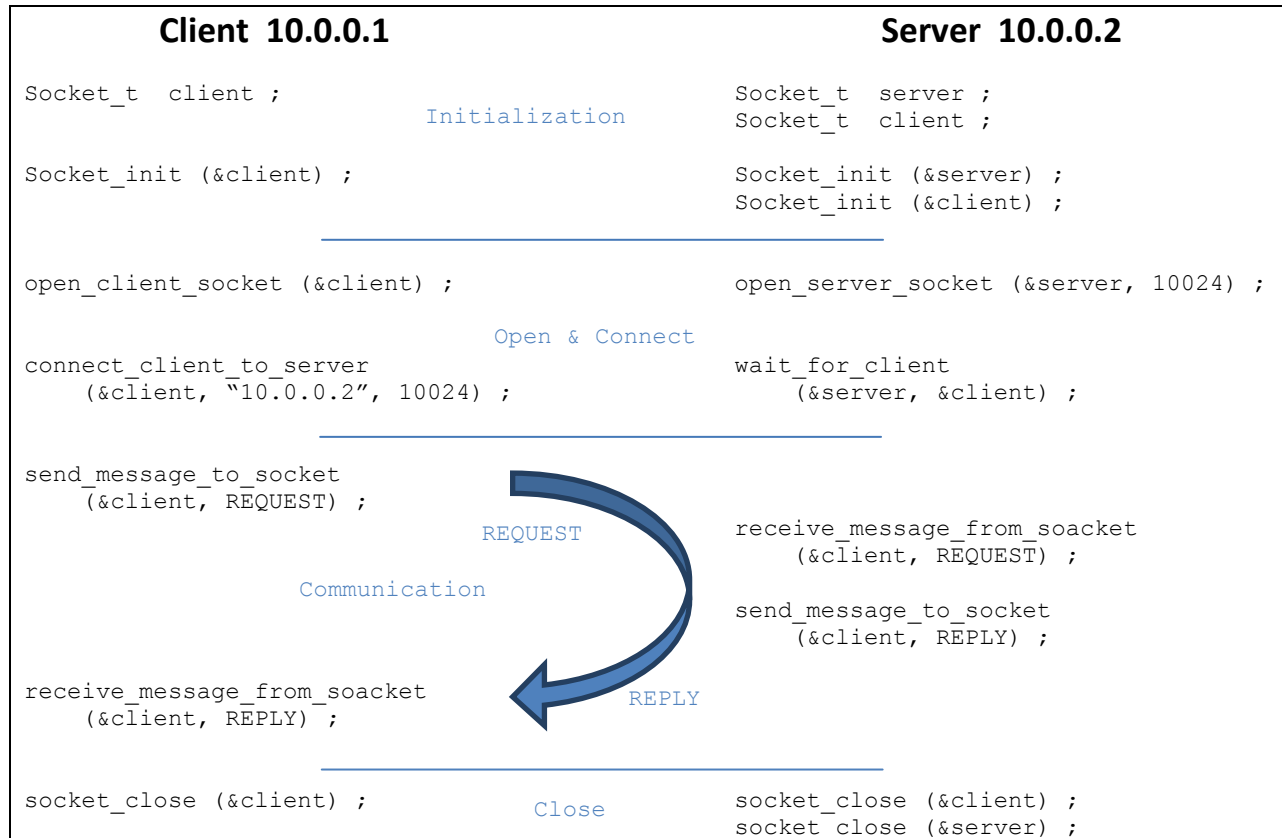
To learn more about all the functions and data structures described above, please refer to the Doxygen documentation of 3D-ICE in the `3D-ICE/doc/html/` folder.

D. Socket

The client / server communication can be implemented using the data type `Socket_t`. To use it, the header file `network_socket.h` must be included in the source file. The documentation of the functions that can be used to init, open and close the network sockets, as well as to send or receive messages are available through the Doxygen documentation.


The following example resumes the communication pattern to follow to establish an exchange of messages between the two processes. Please note that the server must use the data structure of type

Socket twice: one to handle a socket that waits for connections from a client (this is called `server` in the example) and a second one to use to communicate with the client (called `client` in the example) after the connection between the two processes has been established. This mechanism of using two different sockets allows the server to handle many clients at the same time but this functionality is not shown (implemented) in `3D-ICE-Server.c`.



E. NetworkMessage

The data type `NetworkMessage_t`, available including the header file `network_message.h`, stores the message that can be sent over the network sockets to exchange information between two processes. The structure of the messages that can be built with 3D-ICE is the following:



The header of the message contains two words: `Length` and `Type`. `Length` indicates the total length of the message (header plus payload), expressed as number of words that are going to be sent or received through the socket, while `Type` indicates the type of data that the message contains. The messages can have any length and the functions available in `network_message.h` guarantee that the memory to store the messages is handled transparently: the user must not worry about the content of the field `length`

because its value is computed or updated internally whenever new data is added in the payload. The types of network messages supported so far in 3D-ICE are defined through the enumeration `MessageType_t` defined in the header file `types.h`. The user is free to add new message types and to define the corresponding payload as well as the operations that either the client or the server must do whenever they send or receive a message of such type.

The sequence of functions that must be used to create and send a message is the following:

- `network_message_init (NetworkMessage_t *)`
- `build_message_head (NetworkMessage_t *, MessageType_t)`
- `insert_message_word (NetworkMessage_t *, void *)`
- `send_message_to_socket (Socket_t *, NetworkMessage_t *)`
- `network_message_destroy (NetworkMessage_t *)`

The function called to insert a word of data in the payload must be called as many times as the number of words that must be present in the message according to its type. When the message is built, **the payload behaves as a stack**, meaning that every time that a word is inserted with `insert_message_word`, that word will be put at the right of the actual payload and the length of the message will automatically be increased by one unity. Therefore, **the order of the calls** to `insert_message_word` **must reflect the order of the data** in the payload. Because of this mechanism, the function `network_message_destroy` must be called before creating a new message using the same variable to delete the header and the payload and reset the state of the message.

On the other side, when a message must be received from a socket, the sequence of functions to be used is the following:

- `network_message_init (NetworkMessage_t *)`
- `receive_message_from_socket (Socket_t *, NetworkMessage_t *)`
- `extract_message_word (NetworkMessage_t *, void *, int)`
- `network_message_destroy (NetworkMessage_t *)`

Once the message has been received, the function `extract_message_word` is meant to extract a word from the payload using an index. It means that once the message is received, **the words in the payload can be accessed and read in any order**. Again, the number of words in the payload depends on the type of the message that is received (the type can be accessed using the corresponding field in the data structure) and the message must be reset with `free_network_message` before using the same variable to receive a new message.

More details about the usage of the network interface in 3D-ICE can be found in the sources itself or in their documentation. As main reference, the user can follow the two examples available in the `bin` folder: `3D-ICE-Client.c` and `3D-ICE-Server.c`. These two files show how to establish the communication between a client and a server and how a client can build messages to send to the server a sequence of power traces and get back from the server some information about the thermal state of the simulation.

F. Debugging of ThermalSimulation

For the purpose of debugging, several pre-processing options can be enabled to directly check the values computed during the construction of thermal data (thermal grid/circuit, system matrices, sources etc) before the simulation even starts. These options can be activated uncommenting the corresponding line in the file `3d-ice/sources/Makefile` and running the `make` command again (run the `make clean` command before building). As a consequence, messages will be redirected to `stderr`.

The debug options that are available are:

- `PRINT_SYSTEM_MATRIX` prints the content of the system matrix while it is filled. For every column in the system matrix ($\mathbf{G} + \mathbf{C}/h$, see [1] for more details), it prints the cell ID of the corresponding cell and the list of the nonzero coefficients indicating representing its neighbors. For every neighbor, it also prints the cell ID.
- `PRINT_SYSTEM_VECTOR` prints the content of the system vector (the Thermal state) at every time step.

10. References

- [1] A Sridhar, A Vincenzi, M Ruggiero, T Brunschwiler, D Atienza, "3D-ICE: Fast compact transient thermal modeling for 3D-ICs with inter-tier liquid cooling", *Proceedings of the 2010 International Conference on Computer-Aided Design (ICCAD 2010)*, San Jose, CA, USA, November 7-11 2010.
- [2] A Sridhar, A Vincenzi, M Ruggiero, T Brunschwiler, D Atienza, "Compact transient thermal model for 3D ICs with liquid cooling via enhanced heat transfer cavity geometries", *Proceedings of the 16th International Workshop on Thermal Investigations of ICs and Systems (THERMINIC'10)*, Barcelona, Spain, 6-8 October, 2010.
- [3] T Brunschwiler, S Paredes, U Drechsler, B Michel, W Cesar, Y Leblebici, B Wunderle, H Reichl, "Heat-removal performance scaling of interlayer cooled chip stacks", *Proceedings of the 2010 IEEE Intersociety Conference on Thermal and Thermomechanical Phenomena in Electronic Systems (ITHERM '10)*, pp. 1-12, June 2010.
- [4] T A Davis and E P Natarajan, "Algorithm 8xx: KLU, a direct sparse solver for circuit simulation problems," *ACM Transactions on Mathematical Software*, vol.5, no.1, pp.1-14, 2010.
- [5] Doxygen, URL: <http://www.stack.nl/~dimitri/doxygen/>.
- [6] D. Atienza, P. G. Della Valle, G. Paci, F. Poletti and L. Benini, "HW-SW Emulation Framework for Temperature-Aware Design in MPSoCs", *ACM Transactions on Design Automation for Embedded Systems (TODAES)*, vol. 12, no. 3, August, pp. 1 – 26, 2007.
- [7] Lee, S., Song, S., Au, V., Moran, K., "Constriction/Spreading Resistance Model for Electronics Packaging," *ASME/JSME Thermal Engineering Conf.*, Vol.4, 1995, pp.199-206.